

IPv6

Content

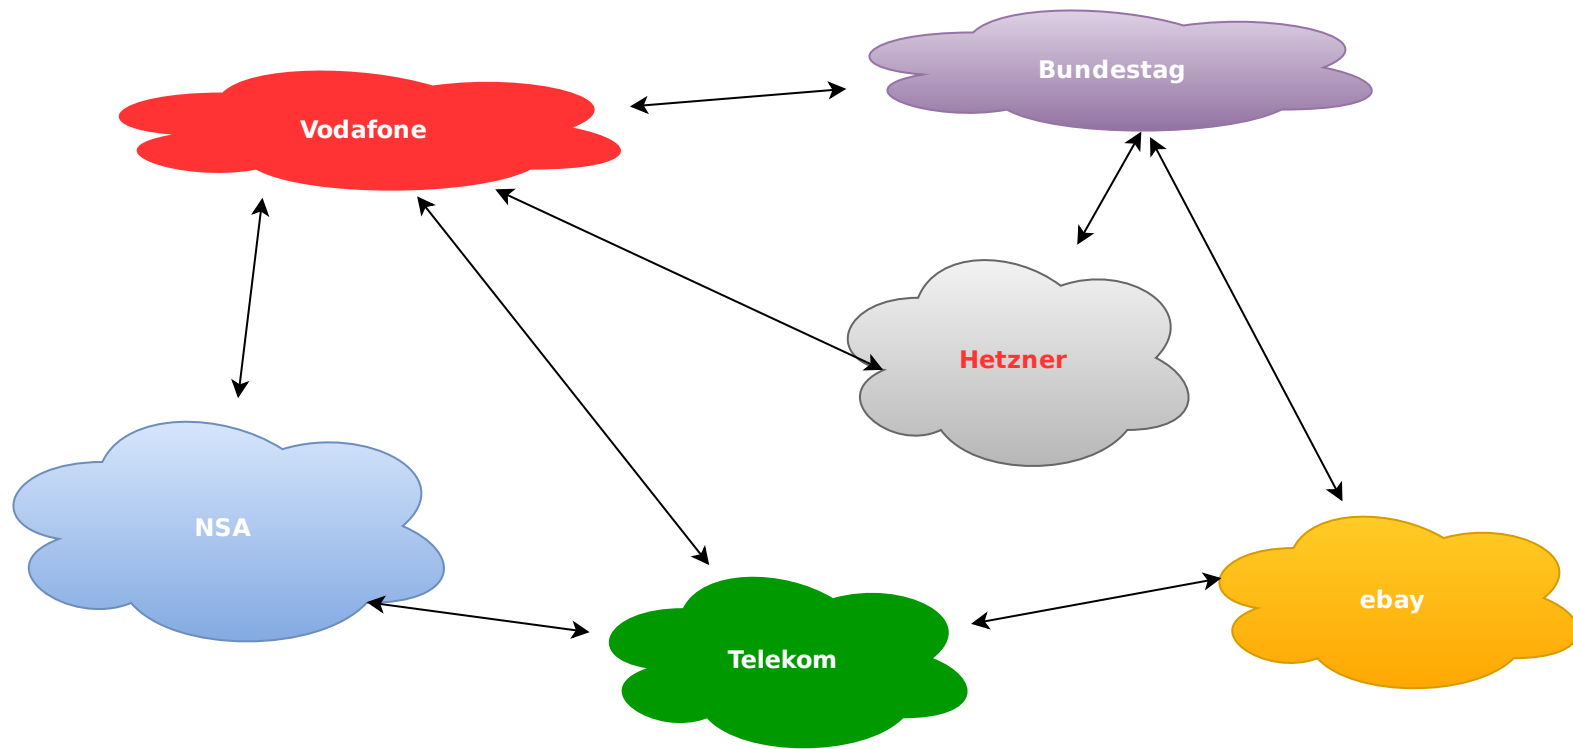
1. Why we need IPv6 (or why IPv4 sucks)
2. New Features
3. Implementation of IPv6

Why do we need IPv6


Why we need IPv6

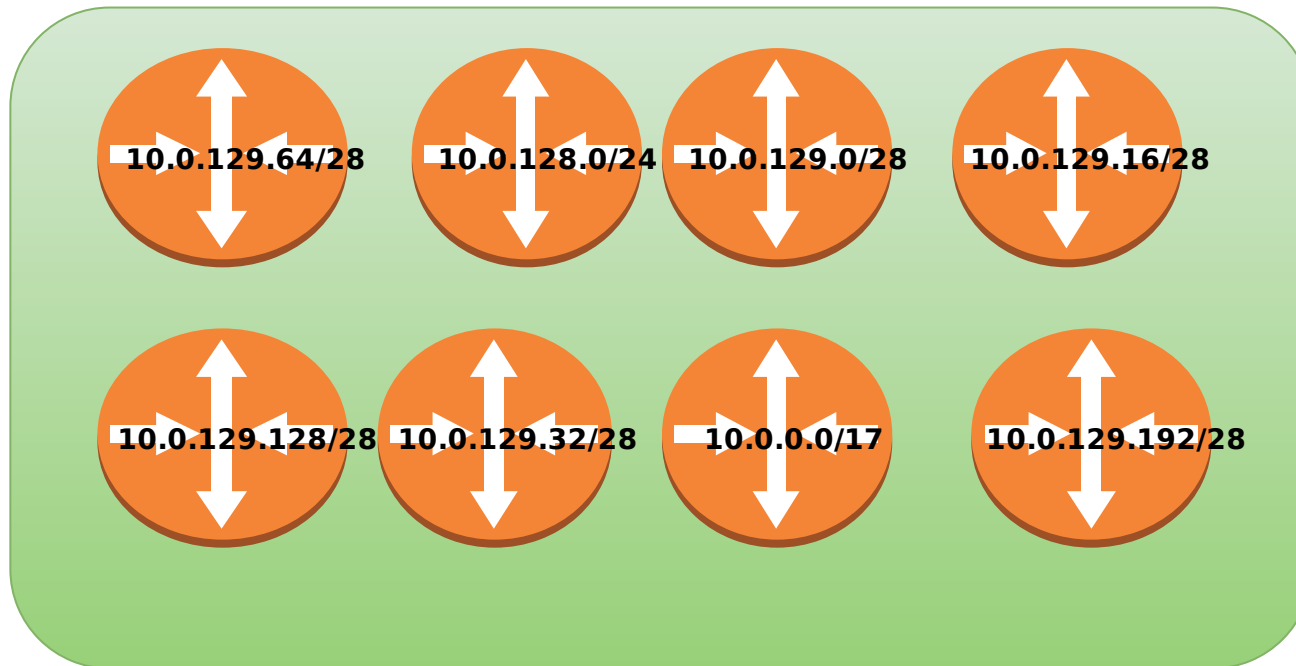
- IPv4 Basics

- 32-bit address range
- 2^{32} addresses => 4 billion
- 1.2 billion active devices
- 10 years ago everybody got a class C or B net
- Less available nets => smaller allocation
- 595,385 BGP entries in full table
- Several router fail at 524288 entries ($512 * 1024$)



Every cloud = One AS (Autonomous





Features

Features

- Moar!! Addresses

- 128-bit length
- 2^{128} addresses
- 7.9×10^{28} times bigger than IPv4 space
- 29-bit prefix for each LIR (Local Internet Registry)
- LIR can get even one per location
- Even 19-bit prefix if you blame the RIR (Regional Internet Registry)
- For example DTAG got 2003::/19
- Every site should get their own 64-bit prefix
- 48-bit minimal PI space
- 2a00::/12 for RIPE
- IANA currently uses 2000::/3

Features

- Moar!!

Addresses

- New notation

- New hex notation: 2a01:4f8:11a:b1f::4/64
- Full notation: 2a01:04f8:011a:0b1f:0000:0000:0000:0004
- 8 groups of 16-bit in hex
- Ignore leading 0s in groups
- Replace the biggest successive group of zeros with ::
- /64 indicates the prefix length
- 64-bit is default, can be omitted

Features

- Moar!!

Addresses

- New notation

- Improvement

- No broadcast
- Multicast is mandatory
- Neighbor Discovery Protocol (NDP) as ARP replacement
- Always assign 64bit prefixes to devices, no single IP-Addresses
- Privacy extension for host part of addresses

Features

- Moar!!

Addresses

- New notation

- Improvement

- Multicast

Excursion about Multicast

- Own address space, prefix is ff::/8
- used for communication in groups
- different groups and scopes available
- Simply assign an address from a group to you to be part of the group
- ff0X::1 : every device with IPv6 (like IPv4 broadcast)
- ff0X::2 : all router
- ff0X::f : UPnP devices
- ff0X::101 : NTP server
- ff0X::1:2 : DHCPv6 server
- Where X is: 1 = Interface-Local; 2 Link-Local, 8 Organization-Local

Features

- Moar!!

Addresses

- New notation

- Improvement

- Multicast

- Useless stuff

- Mobile IPv6
- SLAAC - Stateless Address AutoConfiguration
- DAD - Duplicate Address Detection
- Fake RA - Router Advertisement
- RA Flooding
- Fake DAD

Features

- Moar!!

Addresses

- New notation

- Improvement

- Multicast

- Useless stuff

- Good to know

- IPv6 is always preferred over IPv4
- IPv4 and IPv6 are separate networks, awesome for breaking stuff
- BIOS PXE boot will never support IPv6
- UEFI PXE boot supports IPv6 + IPv4
- Link local address now contains MAC address
fe80::52e5:49ff:fe41:c204
- Route your prefix to server, no bridging
- STOP SHITTY NAT!


Implementation

Implementation

- Global Routing

- IPv6 addresses are 4 times bigger than IPv4
- Requires 4 times more space for each BGP entry
- Routers are already full ...
- Bigger prefixes for each AS -> less prefixes -> less BGP entries
- Hierarchical routing
- Renumber a server if he changes the location, don't route his prefix
- DO NOT ROUTE A 64bit PREFIX

SERIOUSLY, DO NOT ROUTE A SINGLE PREFIX



Implementation

- Global Routing

- System configuration

- Static for servers
- Always try to use link local gateway fe80::1
- SLAAC for clients in trusted networks
- DHCP6 or static for clients in untrusted networks

Implementation

- Global Routing

- System configuration

- Debugging

- Do you run a firewall? Does it care about IPv6?
- You killed IPv4/IPv6 Routing/Firewalling? The other protocol will still work
- Never forget that IPv6 is always preferred over IPv4
- Local issues? Check the cache size of your switch
- Routing issues? Check the cache size of your routers
- Initial connection is slow? Src + Dest have IPv6 configured but it is broken => fallback to IPv4

Summary

Plan your hierarchy wisely

Honor your hierarchy

Be careful with SLAAC